

ALTA 2025

**Proceedings of the 23rd Annual Workshop of the
Australasian Language Technology Association**



November 26-28, 2025
The University of Sydney
Sydney, Australia

The ALTA organizers gratefully acknowledge the support from the following sponsors.

Gold



Silver



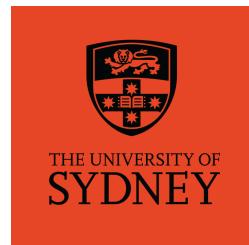
Bronze



GE Aerospace



Host sponsors



Introduction

Welcome to the **23rd Annual Workshop of the Australasian Language Technology Association (ALTA 2025)**. Hosted on the campus of the University of Sydney, ALTA 2025 will provide a platform for the exchange of ideas, exploration of innovations, and discussion of the latest advancements in language technology. The workshop acknowledges the significance of its location on the traditional lands of the Gadigal people of the Eora nation, underscoring a commitment to inclusivity and respect. ALTA 2025 convenes leading researchers, industry experts, and practitioners in the fields of natural language processing (NLP) and computational linguistics.

The program at ALTA this year reflects a wide variety of topics of relevance to contemporary NLP research. This year, ALTA had three options for submission: direct archival papers (long and short), direct non-archival papers (various lengths), and a commitment site for papers reviewed through ACL Rolling Review (ARR). We received 24 archival submissions, accepting 13 after double-anonymous reviews by 3 reviewers per paper. This gives an overall acceptance rate of 54%, which is comparable to previous iterations of ALTA. The papers will be presented in four thematic sessions: natural language understanding, natural language generation, robustness and security of LLMs, and multimodality.

This year, we introduced an alternative to present papers published at international conferences such as ACL, EMNLP, with the view to enable ALTA to be a local forum to disseminate internationally acknowledged research and facilitate discussions among researchers in the region. This was in addition to non-archival papers which may be early work or research proposals. We received 14 submissions in the non-archival track and accepted 13 of them. This includes 8 papers that have been published elsewhere. Every non-archival submission was reviewed by one reviewer. The non-archival papers will be presented as lightning talks across three thematic sessions.

The workshop includes a tutorial on ‘Alignment of Large Language Models with Human Preferences and Values’ by presenters from Macquarie University and a shared task on linking adverse drug events to medical ontologies. This year’s panel discussion is on ‘Teaching/Using NLP for teaching’ and will be moderated by Aditya Joshi. We thank our panelists for offering their insights on the topic. We thank the tutorial presenters, shared task chair and participants. The shared task participants will also present short talks at the workshop.

This year’s submissions showcase advancements across a wide array of topics. The thematic sessions highlight the range of topics covered in the workshop. Scientific literature, code and social media data are recurring text forms. The papers on language and multimodality highlight the narrowing gap between different modalities as foundation models progress towards unified modeling.

We would like to extend our heartfelt gratitude to our keynote speakers: Prof. Alistair Knott - (Victoria University of Wellington), Prof. Mark Johnson (Macquarie University and Oracle). A special thank you to Prof. Massimo Piccardi (University of Technology Sydney) who was the judge for the best paper awards.

Thank you to the members of the organising committee and volunteers for their hard work in preparing and running ALTA. We extend our heartfelt appreciation to the reviewers for their effort in writing detailed and constructive reviews.

Lastly, ALTA 2025 gratefully acknowledges the support of our sponsors: Google (Gold tier), Amazon Nova and Mongo DB (Silver tier), and GE Aerospace and University of Melbourne (Bronze tier). We are also proud to have the University of Sydney as our host.

The success of this workshop would not be possible without your invaluable contributions. Welcome to USyd and Sydney! We hope that you enjoy ALTA 2025, and look forward to a rewarding and inspiring time together.

Jonathan K. Kummerfeld
ALTA 2025 General Chair

Aditya Joshi
Mark Dras
ALTA 2025 Program Chairs

Organizing Committee

General Chair

Jonathan K. Kummerfeld, The University of Sydney

Program Chairs

Aditya Joshi, University of New South Wales
Mark Dras, Macquarie University

Publication Chair

Ana Clarissa Miranda Pena, The University of Sydney

Technology Chair

Ngoc Gia Hy Nguyen, The University of Sydney

Sponsorship Chair

Yidong Gan, The University of Sydney

Local Chairs

Sebastian Kobler, The University of Sydney
Thomas Elton, The University of Sydney

Program Committee

Reviewers

Naomi Baes, University of Melbourne
Daniel Beck, Royal Melbourne Institute of Technology
Dawei Chen, Australian National University
Ming-Bin Chen, University of Melbourne
Xiang Dai, CSIRO
Mark Dras, Macquarie University
Gabriela Ferraro, Australian National University
Eduard Hovy, University of Melbourne
Nicholas I-Hsien Kuo, University of New South Wales
Inigo Jauregi Unanue, University of Technology Sydney
Antonio Jimeno Yepes, Royal Melbourne Institute of Technology
Jonathan K. Kummerfeld, University of Sydney
Sarvnaz Karimi, CSIRO
Kemal Kurniawan, University of Melbourne
Chunhua Liu, University of Melbourne
Raphael Merx, University of Melbourne
Diego Mollá, Macquarie University
Massimo Piccardi, University of Technology Sydney
Josiah Poon, University of Sydney
Anudeex Shetty, University of Melbourne
Amrita Singh, University of New South Wales
Dipankar Srirag, University of New South Wales
Hanna Suominen, Australian National University
Zheng Wei Lim, University of Melbourne
Qiongkai Xu, Macquarie University

Keynote Talk: How can computational linguists contribute to the social good?

Prof. Alistair Knott

Victoria University of Wellington

2025-11-27 10:00:00 – Room: Lecture Theatre 302, PNR building

Abstract: Like most AI language researchers these days, I divide my time between being amazed by the quality of our new large language models, and worrying about how these models will impact society. In an attempt to reconcile these two instincts, my work for the last few years has focussed on contributing to conversations about LLM regulation, where people who understand the technology at issue have many important roles to play.

In this talk I'll discuss a few ongoing conversations I'm involved in - and hopefully connect with people at ALTA who are having similar conversations. One discussion is about transparency for AI-generated content. My group at the Global Partnership on AI helped to define a clause in the EU's AI Act, requiring providers of AI content generators to ensure their content is 'detectable'. A key debate at present is about enforcement: what can we require of providers, given the state-of-the-art in AI content detection schemes? Another discussion is about transparency of social media platforms. The EU's Digital Services Act contains provisions allowing vetted external researchers to access the largest platforms, to study 'societal risks' (and their mitigation). Now these provisions are in place, what are the most important questions that should be asked, and what methods are most suitable? (I co-founded a grouping of social data science researchers, to help answer that question.) A final discussion is about accountability for AI training sets. I am particularly interested in the labelled datasets created by Silicon Valley companies, to align their LLMs, and to train their harmful content classifiers. A great deal of power resides with the annotators who provide labels for these datasets. But we often know little about how annotators are chosen, and about the guidelines they follow. My group at GPAI has proposed that labels could be provided 'democratically', by consulting a representative sample of citizens or users, through a process akin to opinion polling. I'll finish by discussing how this idea could feature in the governance of 'sovereign LLMs' developed in Australia and New Zealand.

Bio: Ali Knott is Professor of AI at Victoria University of Wellington. He studied Philosophy and Psychology at Oxford, and did his postgrad and postdoc work at in Edinburgh, in computational linguistics. He moved to New Zealand in 1998, and participated in ALTA during its heady early years.

Ali spent many years researching how language is implemented in the human brain. His main interest was in how language interfaces to the sensory and motor systems, to enable us to talk about what we see and do. He presented his theory about this interface in a book published by MIT Press. He extended the theory in the New Zealand AI company Soul Machines. One application was an embodied model of an 18-month-old toddler, 'BabyX'.

For the last 10 years, Ali's main focus has been on the social impacts and governance of AI. He is involved in many international AI policy discussions, mainly through his work as co-lead of a project on Social Media Governance, coordinated by the Global Partnership on AI (now part of the OECD). This work has had several impacts on EU tech legislation, including on provisions about AI content detection (in the AI Act), and provisions to allow vetted researchers access to the largest online platforms (in the Digital Services Act). Ali also participated in the Christchurch Call to eliminate Terrorist and Violent Extremist Content Online, and contributes to the Forum for Information and Democracy. He has advised the New Zealand government on many questions of AI policy, and often looks across the ditch to see how policy is progressing in Australia.

Keynote Talk: How LLMs Change NLP

Prof. Mark Johnson

Macquarie University and Oracle

2025-11-28 14:00:00 – Room: Lecture Theatre 302, PNR building

Abstract: LLMs have made large-scale industrial application of NLP easier than ever before. This talk surveys the strengths and weaknesses of LLMs as qualitative rather than quantitative models, and focuses on the critical role of training data and post-training. Rather than focusing on superficial linguistic form, NLP development now involves understanding information sources and crafting the information flow between agents. Evaluation and testing are emerging as key challenges for LLM-powered applications.

Bio: Mark Johnson is a Professor of Language Science (CORE) in the School of Computing at Macquarie University. He is also the Chief AI Scientist, Oracle Digital Assistant at Oracle Corporation, where he develops chatbots and digital assistants. The Oracle Digital Assistant division develops novel deep learning models to power the next generation of Conversational AI using semantic parsing.

Mark Johnson has worked on a wide range of topics in computational linguistics, but his main area of research is natural language understanding, especially syntactic parsing and semantic analysis, and their applications to text and speech processing.

Table of Contents

<i>Robustness of Neurosymbolic Reasoners on First-Order Logic Problems</i>	
Hannah Bansal, Kemal Kurniawan and Lea Frermann	1
<i>Nek Minit: Harnessing Pragmatic Metacognitive Prompting for Explainable Sarcasm Detection of Australian and Indian English</i>	
Ishmanbir Singh, Dipankar Srirag and Aditya Joshi	13
<i>Some Odd Adversarial Perturbations and the Notion of Adversarial Closeness</i>	
Shakila Mahjabin Tonni, Pedro Faustini and Mark Dras	28
<i>Thinker-DDM: Modeling Deliberation for Machine Translation with a Drift-Diffusion Process</i>	
Hongbin Na, Zimu Wang, Mieradilijiang Maimaiti, Tong Chen, Wei Wang, Tao Shen and Ling Chen	45
<i>Can an LLM Elicit Information from Users in Simple Optimization Modelling Dialogues?</i>	
Yelaman Abdullin, Diego Mollá, Bahadorreza Ofoghi, Vicky Mak-Hau and John Yearwood . .	64
<i>SHIELD: Classifier-Guided Prompting for Robust and Safer LVLMs</i>	
Juan Ren, Mark Dras and Usman Naseem	76
<i>Understanding Multilingual ASR Systems: The Role of Language Families and Typological Features in Seamless and Whisper</i>	
Simon Gonzalez, Tao Hoang, Maria Myung-Hee Kim, Bradley Donnelly, Jennifer Biggs and Tim Cawley	90
<i>An LLM-based Framework for Domain-Specific Information Extraction: A Case Study in Computer Science and Chemistry</i>	
Xungang Gu, Yangjie Tian, Ning Li, Meng Liu, Ruohua Xu, He Zhang, Hanqiu Liu, Yongpan Sheng and Ming Liu	101
<i>Simple and Effective Baselines for Code Summarisation Evaluation</i>	
Jade Robinson and Jonathan K. Kummerfeld	112
<i>MAPLE: Multi-Agent Adaptive Planning with Long-Term Memory for Table Reasoning</i>	
Ye Bai, Minghan Wang and Thuy-Trang Vu	132
<i>A Dataset and Benchmark on Extraction of Novel Concepts on Trust in AI from Scientific Literature</i>	
Melanie McGrath, Harrison Bailey, Necva Böltüçü, Xiang Dai, Sarvnaz Karimi, Andreas Duenser and Cécile Paris	159
<i>LLMs for Argument Mining: Detection, Extraction, and Relationship Classification of pre-defined Arguments in Online Comments</i>	
Matteo Guida, Yulia Otmakhova, Eduard Hovy and Lea Frermann	176
<i>Graph-Score: A Graph-grounded Metric for Audio Captioning</i>	
Manh Luong, Gholamreza Haffari, Dinh Phung and Lizhen Qu	192
Shared Task (Not Peer Reviewed)	
<i>Overview of the 2024 ALTA Shared Task: Normalise Adverse Drug Events</i>	
Diego Mollá, Xiang Dai, Sarvnaz Karimi and Cécile Paris	202

<i>Team MonoLink at the ALTA Shared Task 2025: Synonym-Aware Retrieval with Guideline-Aware Re-Ranking for MedDRA Normalization</i>	
James C. Douglas	208
<i>A Hybrid System for Comprehensive and Consistent Automated MedDRA Coding of Adverse Drug Event</i>	
Abir Naskar, Liuliu Chen, Jemina Kang and Mike Conway	216
<i>SCaLER@ALTA 2025: Hybrid and Bi-Encoder Approaches for Adverse Drug Event Mention Normalization</i>	
Shelke Akshay Babasaheb and Anand Kumar Madasamy.....	224
<i>DRAGON: Dual-Encoder Retrieval with Guided Ontology Reasoning for Medical Normalization</i>	
Dao Sy Duy Minh, Nguyen Lam Phu Quy, Pham Phu Hoa, Tran Chi Nguyen, Huynh Trung Kiet and Truong Bao Tran	230
<i>A Hybrid Retrieval System for Adverse Event Concept Normalization Integrating Contextual Scoring, Lexical Augmentation, and Semantic Fine-Tuning</i>	
Saipriya Dipika Vaidyanathan	240
Tutorial (Not Peer Reviewed)	
<i>Alignment of Large Language Models with Human Preferences and Values</i>	
Usman Naseem, Gautam Siddharth Kashyap, Kaixuan Ren, Yiran Zhang, Utsav Maskey, Juan Ren and Afrozah Nadeem	245

Program

Wednesday, November 26, 2025

14:00 - 17:00 *Tutorial*

17:00 - 18:30 *Meet-n-Greet*

Thursday, November 27, 2025

- 09:10 - 09:20 *Opening*
- 09:20 - 09:30 *Programme Committee Summary*
- 09:30 - 10:00 *Lightning Talks, session 1: Benchmarks & Evaluation*
- 10:00 - 11:00 *ALTA Keynote 1: “How can computational linguists contribute to the social good?” (Professor Alistair Knott).*
- 11:00 - 11:30 *Morning Tea*
- 11:30 - 12:30 *Oral presentations, session 2: Robustness and Safety of Language Models*
- 12:30 - 14:00 *Lunch*
- 14:00 - 15:00 *Panel Discussion [Panellists: Jey Han Lau, Massimo Piccardi, Thomas Elton]*
- 15:00 - 15:30 *Afternoon Tea*
- 15:30 - 16:45 *Oral presentations, session 3: Multimodality and Language*
- 16:45 - 17:15 *Lightning Talks, session 4: Hallucination & Vulnerabilities*
- 17:30 - 21:00 *Social event at Friend in Hand Hotel*

Friday, November 28, 2025

- 09:00 - 09:30 *Lightning Talks, session 5: Model Architecture & Learning*
- 09:30 - 10:30 *Oral presentations, session 6: Language Generation*
- 10:30 - 11:00 *Morning Tea*
- 11:00 - 11:30 *Oral presentations, session 7: Language Understanding*
- 11:30 - 12:30 *Oral presentations, session 8: Shared Task*
- 12:30 - 14:00 *Lunch*
- 14:00 - 15:00 *ALTA Keynote 2: “How LLMs Change NLP” (Professor Mark Johnson).*
- 15:00 - 15:30 *Afternoon Tea*
- 15:30 - 16:00 *ALTA AGM*
- 16:00 - 17:00 *Best Paper Award / Shared Task Award / Closing*